

Kernel versus end-to-end attention speedup

